

Introduction to Big Data with Apache Spark

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



BerkeleyX

This Lecture

The Big Data Problem

Hardware for Big Data **Assignment Project Exam Help**

Distributing Work <https://powcoder.com>

Handling Failures and Slow Machines **Add WeChat powcoder**

Map Reduce and Complex Jobs

Apache Spark

Some Traditional Analysis Tools

- Unix shell commands, Pandas, R

Assignment Project Exam Help

<https://powcoder.com>

All run on a
single machine!
Add WeChat powcoder

The Big Data Problem

- Data growing faster than computation speeds
Assignment Project Exam Help
- Growing data sources
 - » Web, mobile, scientific <https://powcoder.com>
- Storage getting cheaper
Add WeChat powcoder
 - » Size doubling every 18 months
- But, stalling CPU speeds and storage bottlenecks



Big Data Examples

- Facebook's daily logs: 60 TB
- 1,000 genomes project: 200 TB
 - Assignment Project Exam Help
 - <https://powcoder.com>
- Google web index: 10+ PB
 - Add WeChat powcoder
- Cost of 1 TB of disk: ~\$35
- Time to read 1 TB from disk: 3 hours
 - (100 MB/s)

The Big Data Problem

- A single machine can no longer process or even store all the data! **Assignment Project Exam Help**
- Only solution is to ~~distribute~~ data over large clusters
<https://powcoder.com>

Add WeChat powcoder



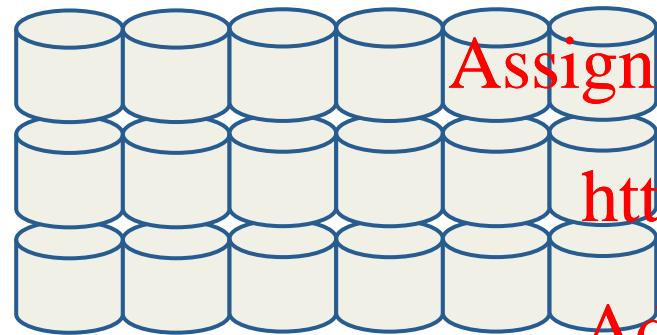
Google Datacenter

Assignment Project Exam Help

How do we program this thing?
<https://powcoder.com>

Add WeChat powcoder

Hardware for Big Data



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Lots of hard drives ... and CPUs



Hardware for Big Data

One big box?

(1990's solution)



<https://powcoder.com>

Add WeChat powcoder

But, expensive

- » Low volume
- » All “premium” hardware

And, still not *big* enough!

Image: Wikimedia Commons / User:Tonusamuel

Hardware for Big Data

Consumer-grade hardware

Not “gold plated”

<https://powcoder.com>

Many desktop-like servers

Easy to add capacity

Add WeChat powcoder

Cheaper per CPU/disk

Complexity in software

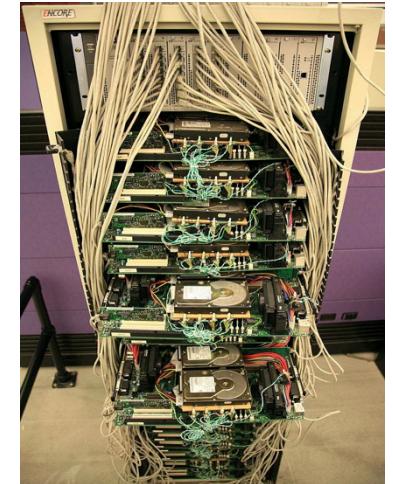


Image: Steve Jurvetson/Flickr

Problems with Cheap Hardware

Failures, Google's numbers:

1-5% hard drives/year

0.2% DIMMs/year

Assignment Project Exam Help

<https://powcoder.com>

Network speeds versus shared memory

Much more latency

Network slower than storage

Uneven performance

What's Hard About Cluster Computing?

- How do we split work across machines?

Assignment Project Exam Help

<https://powcoder.com>

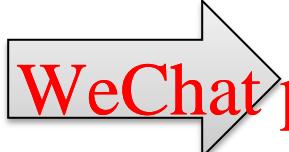
Add WeChat powcoder

How do you count the number of occurrences of each word in a document?

Assignment Project Exam Help

“I am Sam
I am Sam
Sam I am
Do you like
Green eggs and ham?”

I: 3
<https://powcoder.com>
am: 3
Sam: 3
do: 1

Add  WeChat powcoder

you: 1
like: 1

...

One Approach: Use a Hash Table

“I am Sam
Assignment Project Exam Help
I am Sam <https://powcoder.com>
Sam I am [Add WeChat powcoder](#)
Do you like
Green eggs and ham?”

One Approach: Use a Hash Table

“I am Sam Assignment Project Exam Help { | : 1 }

I am Sam <https://powcoder.com>

Sam I am Add WeChat powcoder

Do you like

Green eggs and ham?”

One Approach: Use a Hash Table

“I am Sam Assignment Project {!;.1, Exam Help
I am Sam https://powcoder.com
Sam I am Add WeChat powcoder
Do you like
Green eggs and ham?”

One Approach: Use a Hash Table

“I am **Sam**”
Assignment Project Exam Help
I am Sam <https://powcoder.com>
Sam I am Add WeChat powcoder
Do you like Sam: 1 {
am: 1,

Green eggs and ham?”

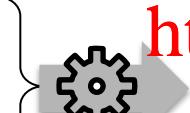
One Approach: Use a Hash Table

“I am Sam Assignment Project Exam Help
I am Sam https://powcoder.com
Sam I am Add WeChat powcoder
Do you like Sam: 1 {
am: 1,
Sam: 1}

Green eggs and ham?”

What if the Document is Really Big?

"I am Sam
I am Sam
Sam I am
Do you like
Green eggs and ham?
I do not like them
Sam I am
I do not like
Green eggs and ham
Would you like them
Here or there?
..."

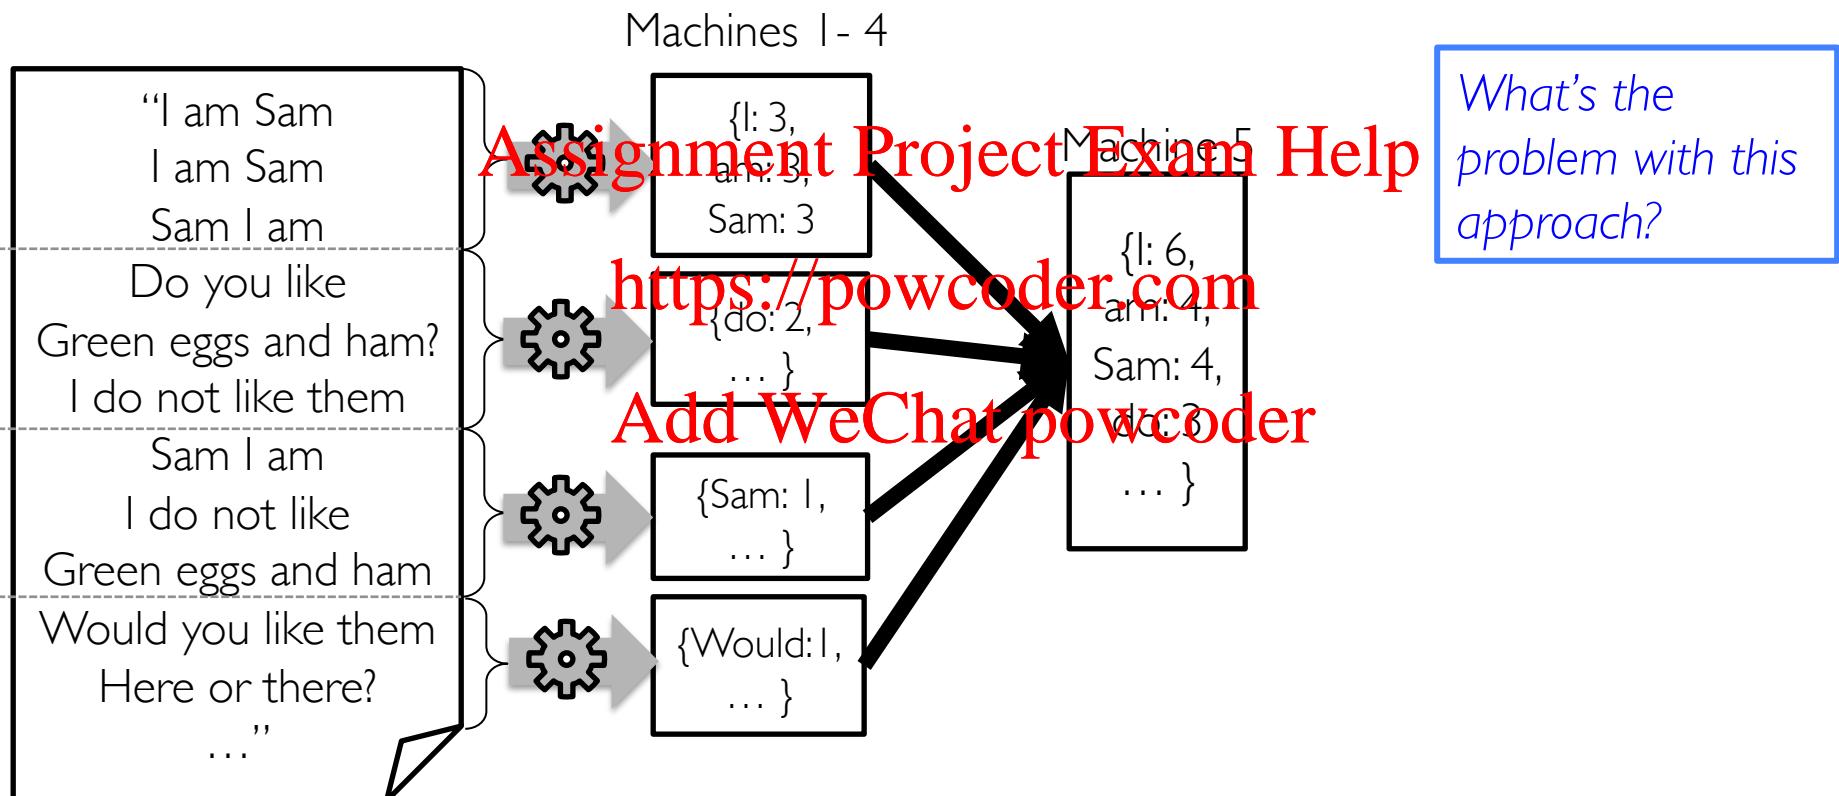


Assignment Project Exam Help

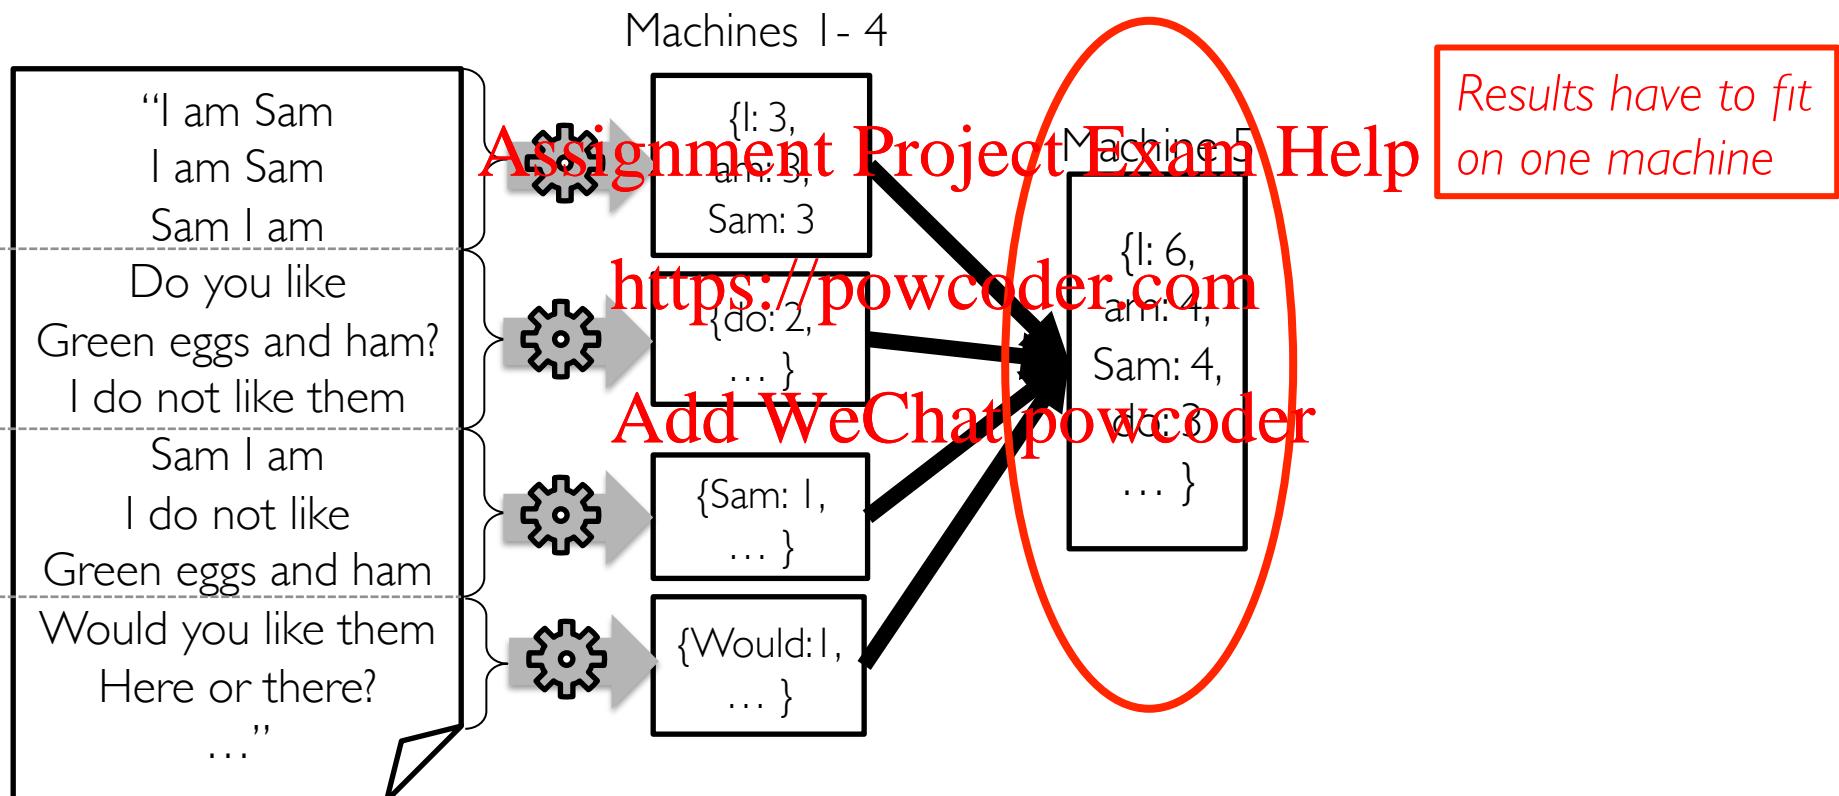
<https://powcoder.com>

Add WeChat powcoder

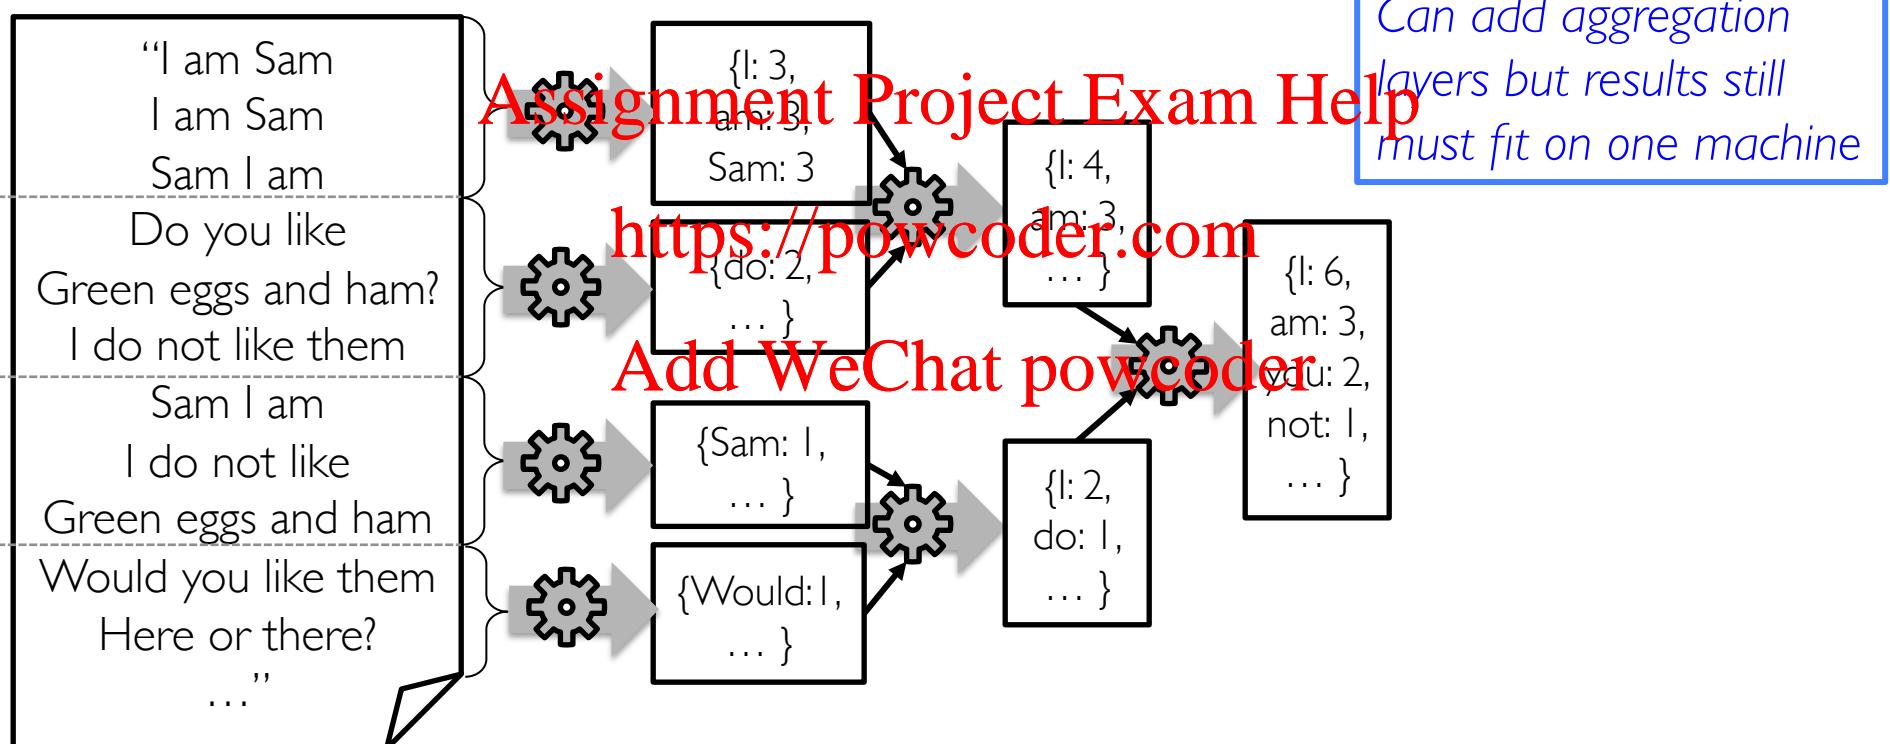
What if the Document is Really Big?



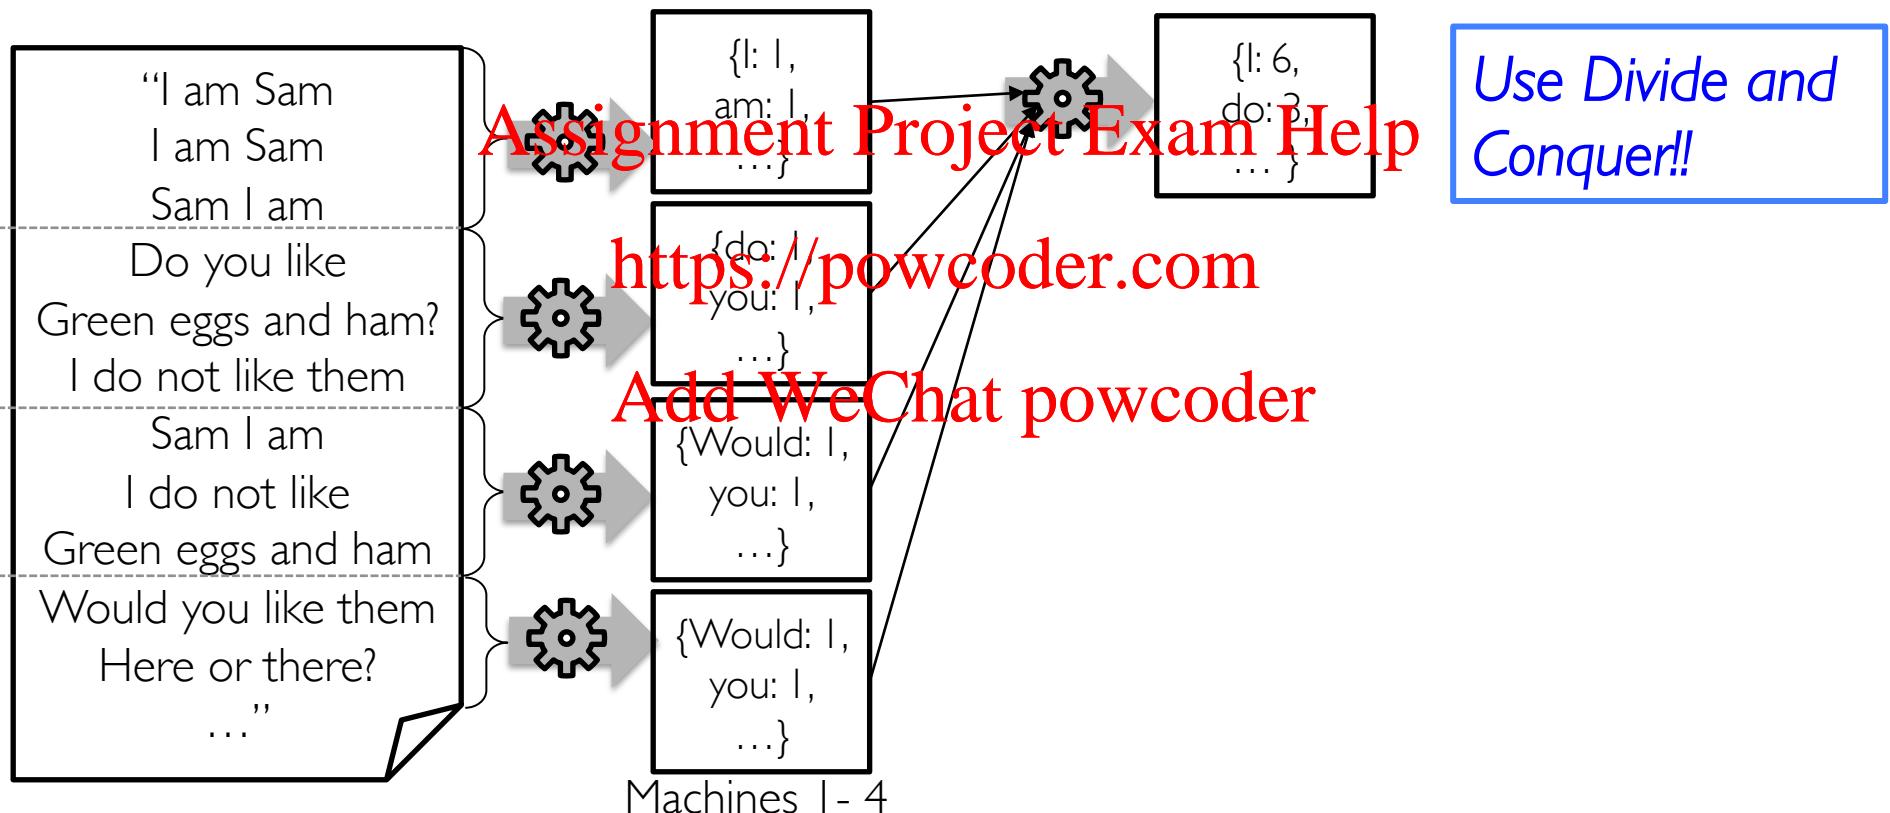
What if the Document is Really Big?



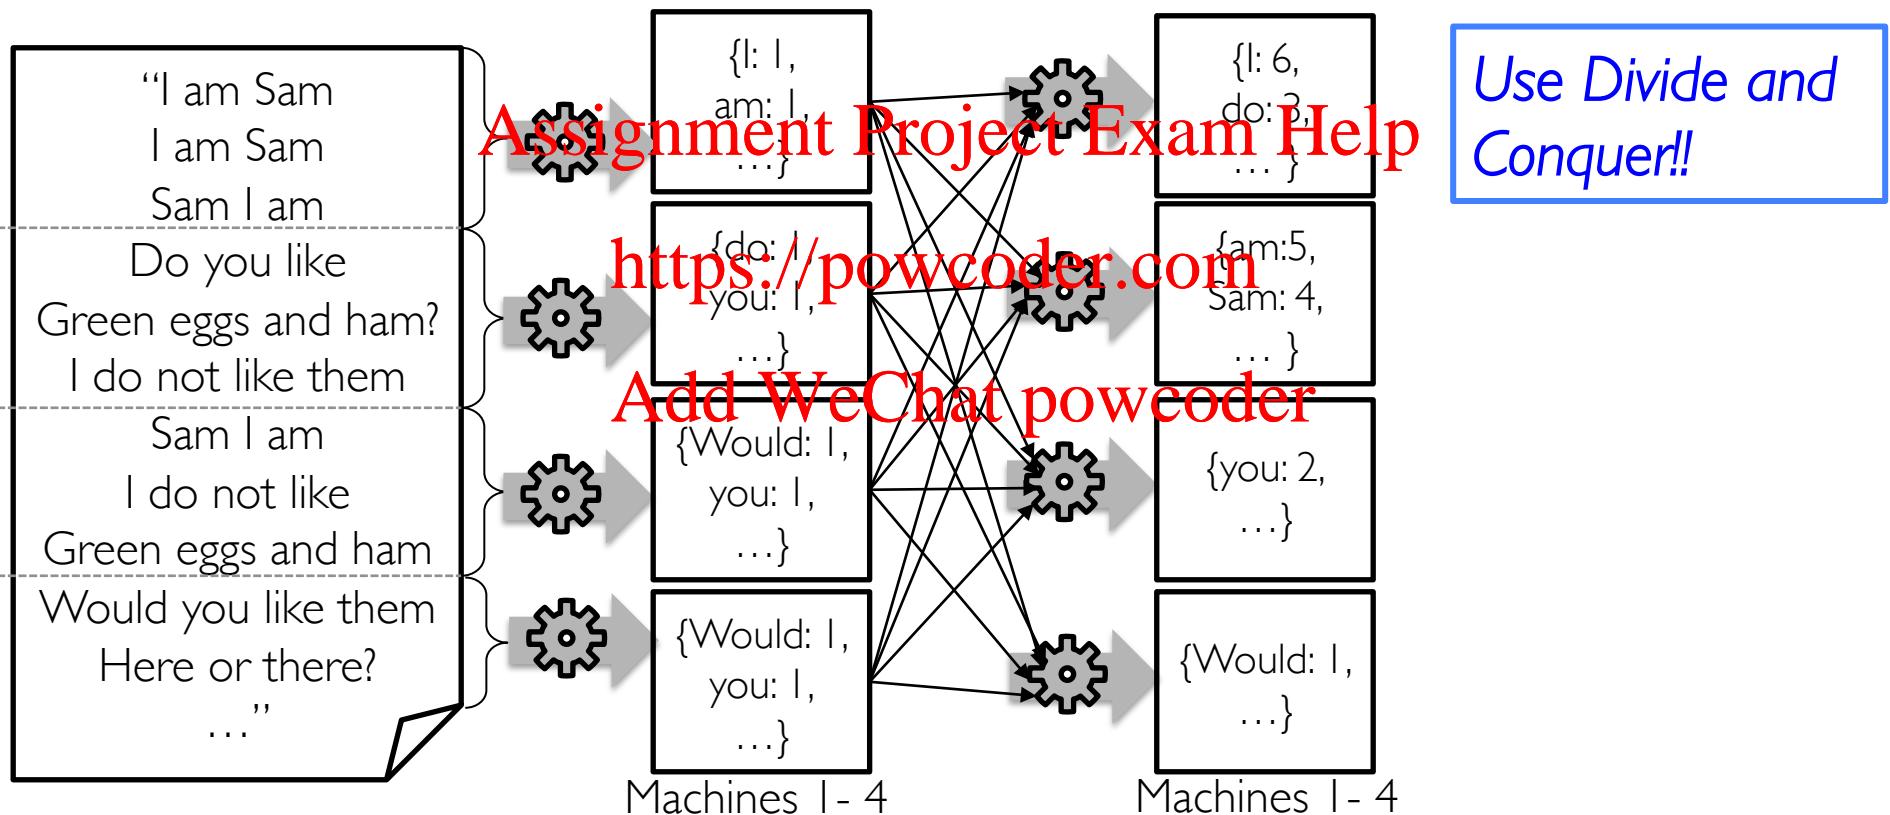
What if the Document is Really Big?



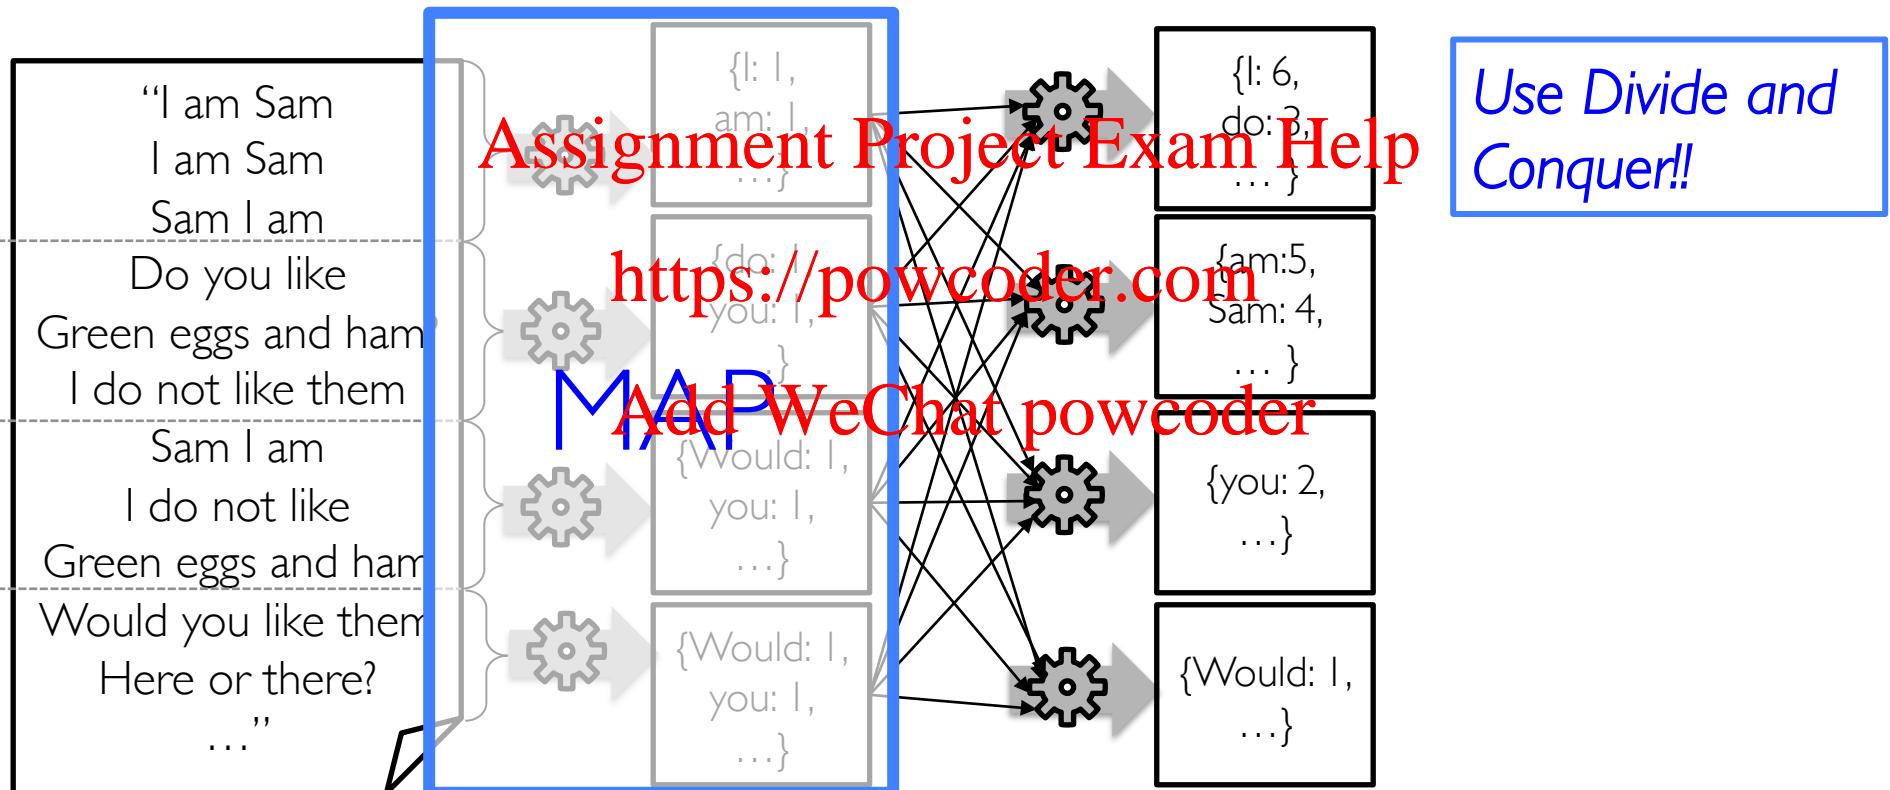
What if the Document is Really Big?



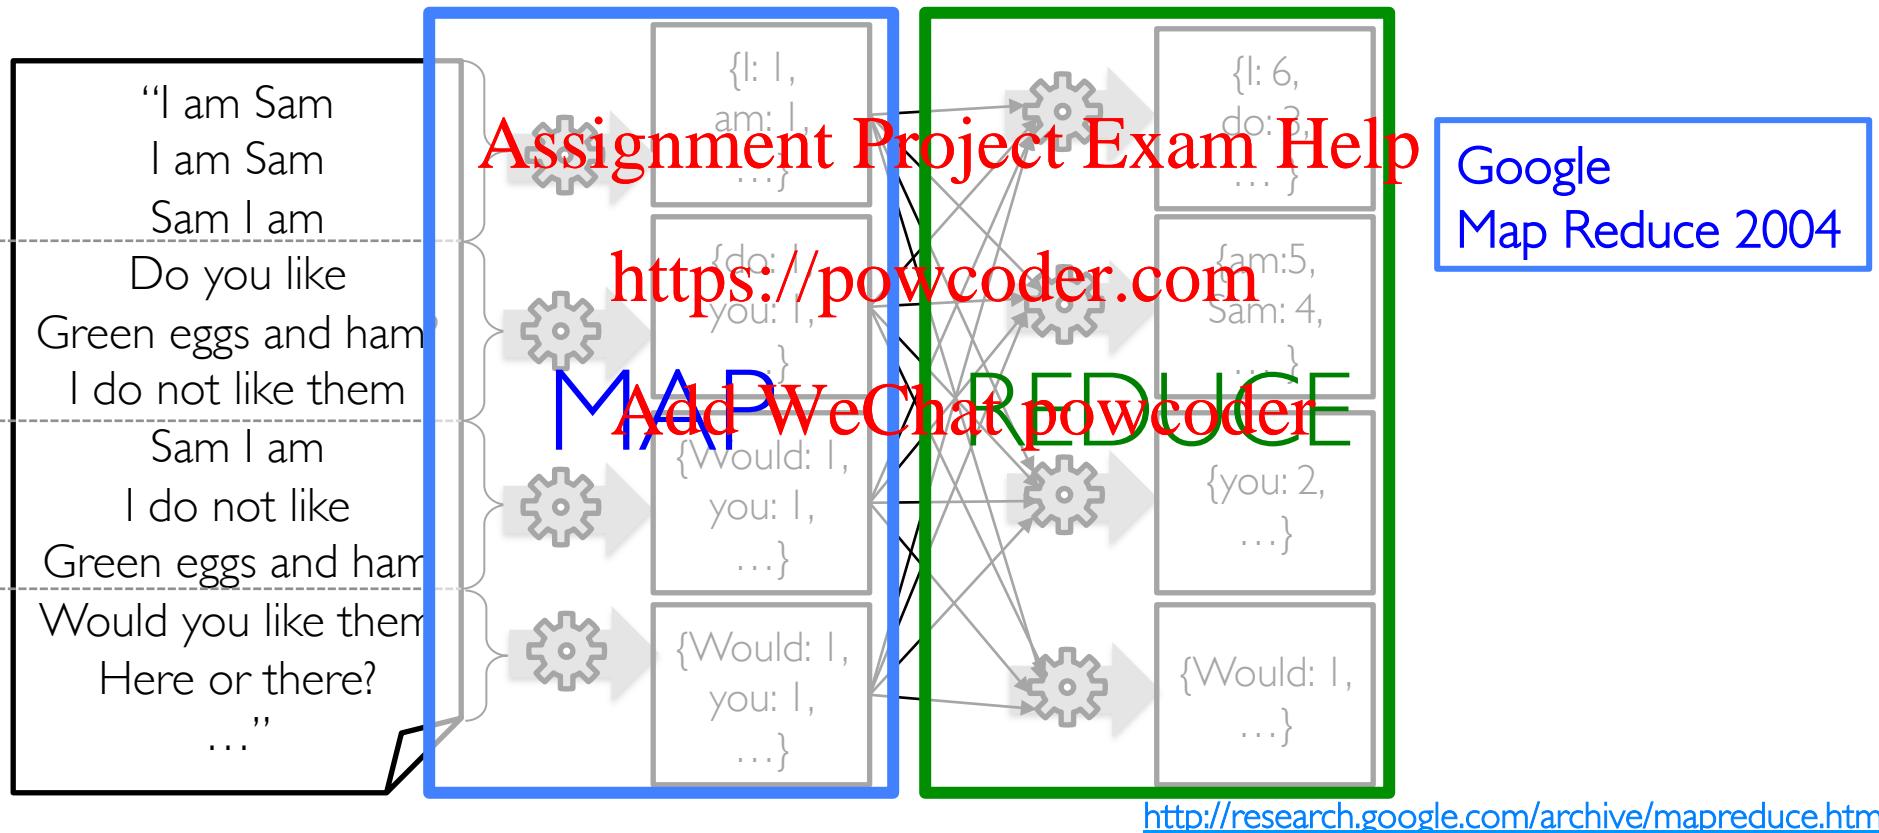
What if the Document is Really Big?



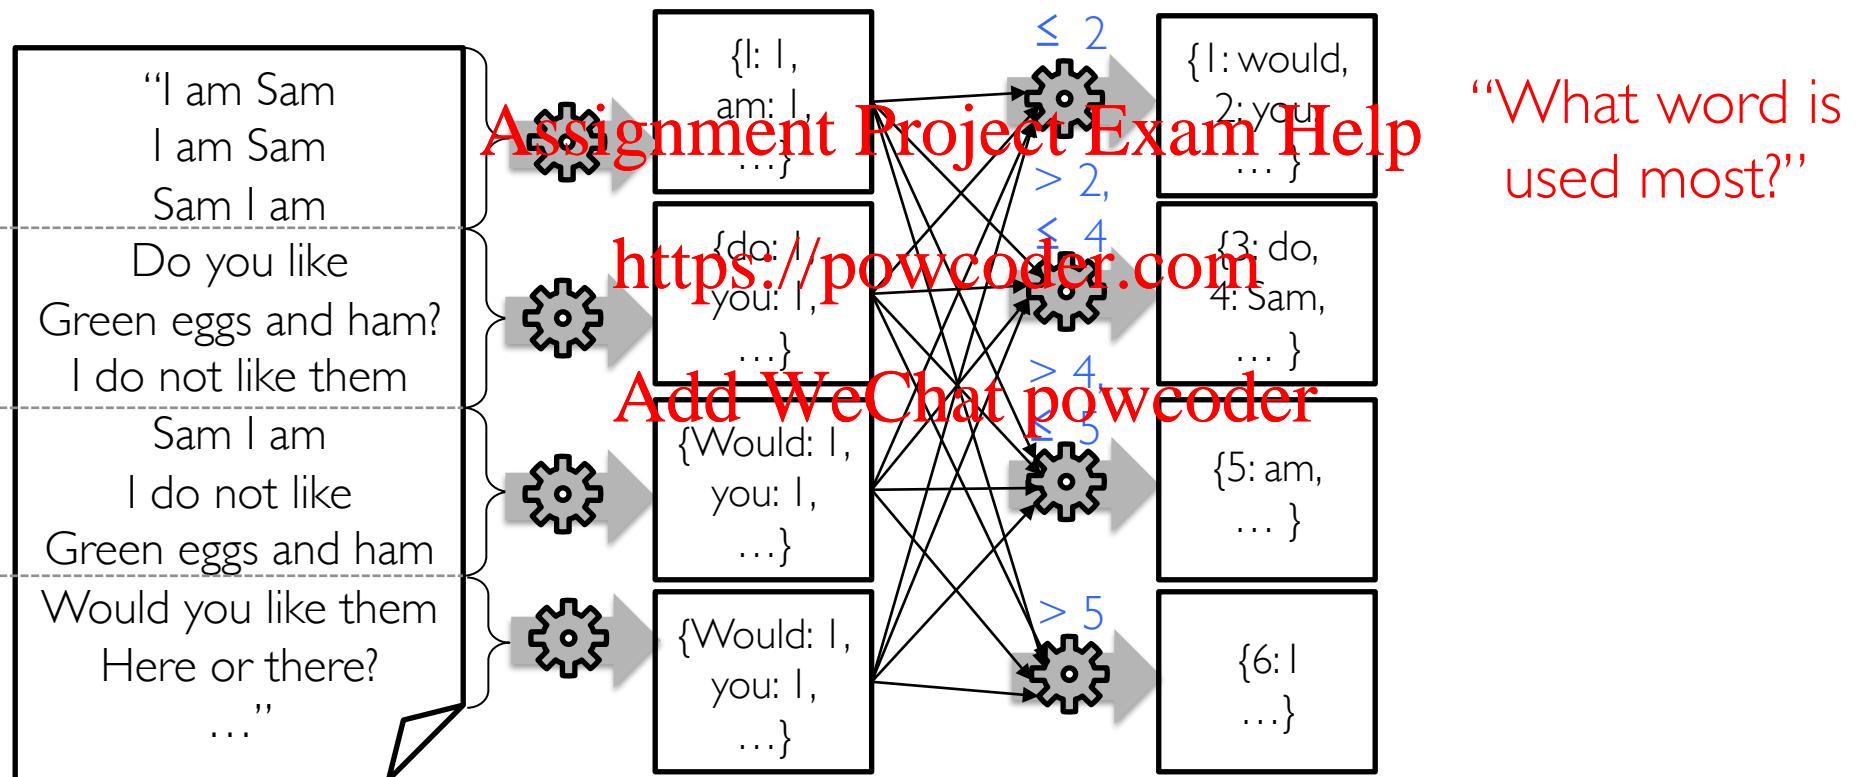
What if the Document is Really Big?



What if the Document is Really Big?



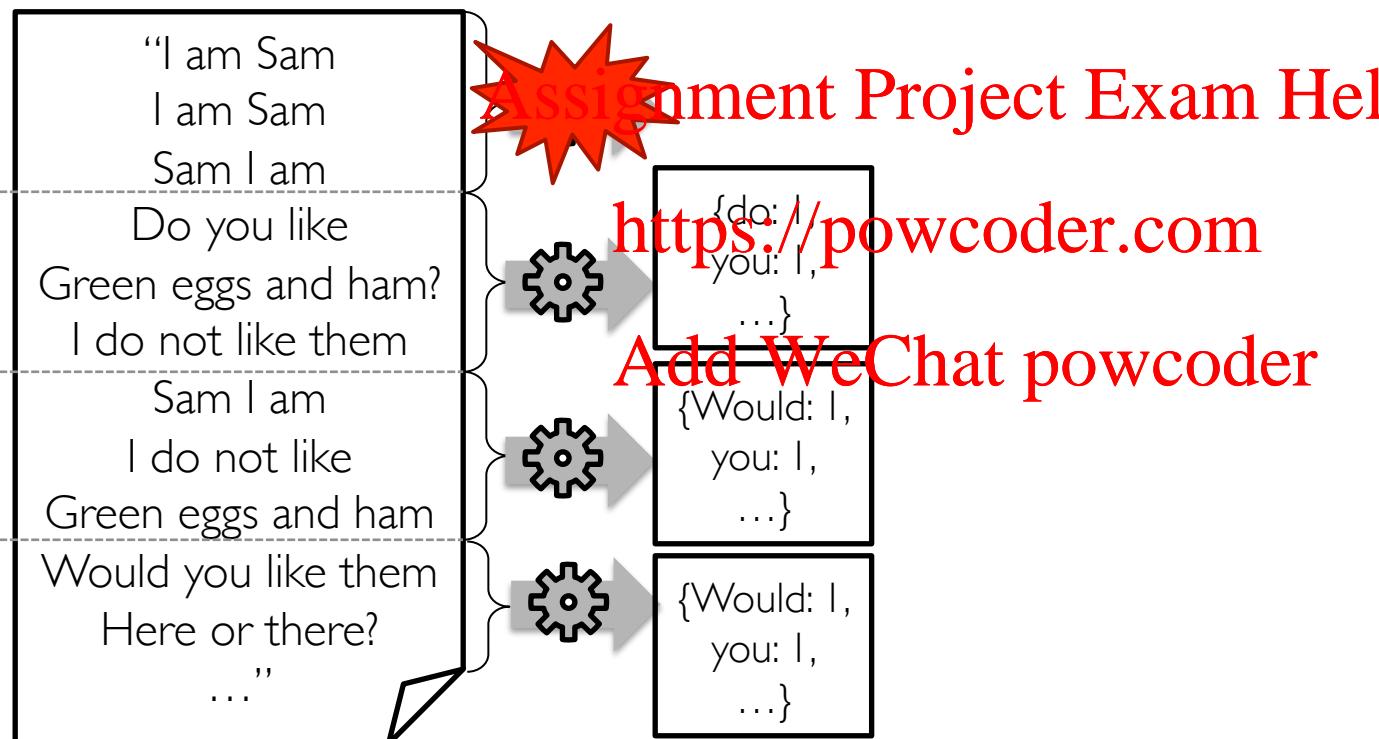
Map Reduce for Sorting



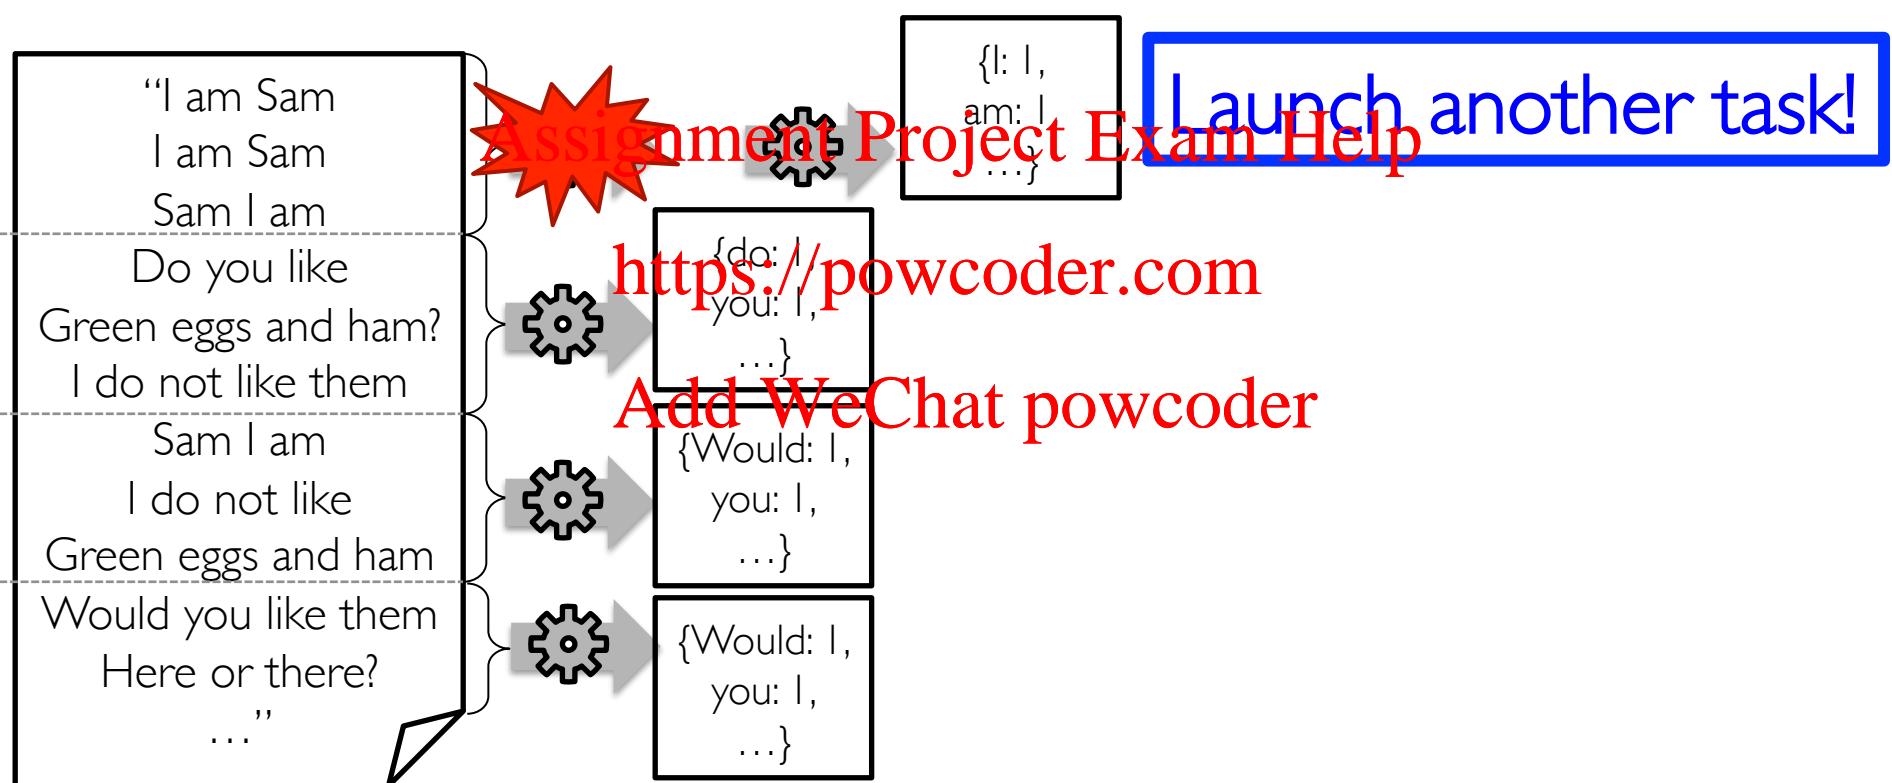
What's Hard About Cluster Computing?

- How to divide work across machines?
 - » Must consider network, data locality
 - » Moving data may be very expensive
- How to deal with failures?
 - » 1 server fails every 3 years → with 10,000 nodes see 10 faults/day
 - » Even worse: stragglers (not failed, but slow nodes)

How Do We Deal with Failures?



How Do We Deal with Machine Failures?



How Do We Deal with Slow Tasks?

"I am Sam
I am Sam
Sam I am
Do you like
Green eggs and ham?
I do not like them
Sam I am
I do not like
Green eggs and ham
Would you like them
Here or there?
..."

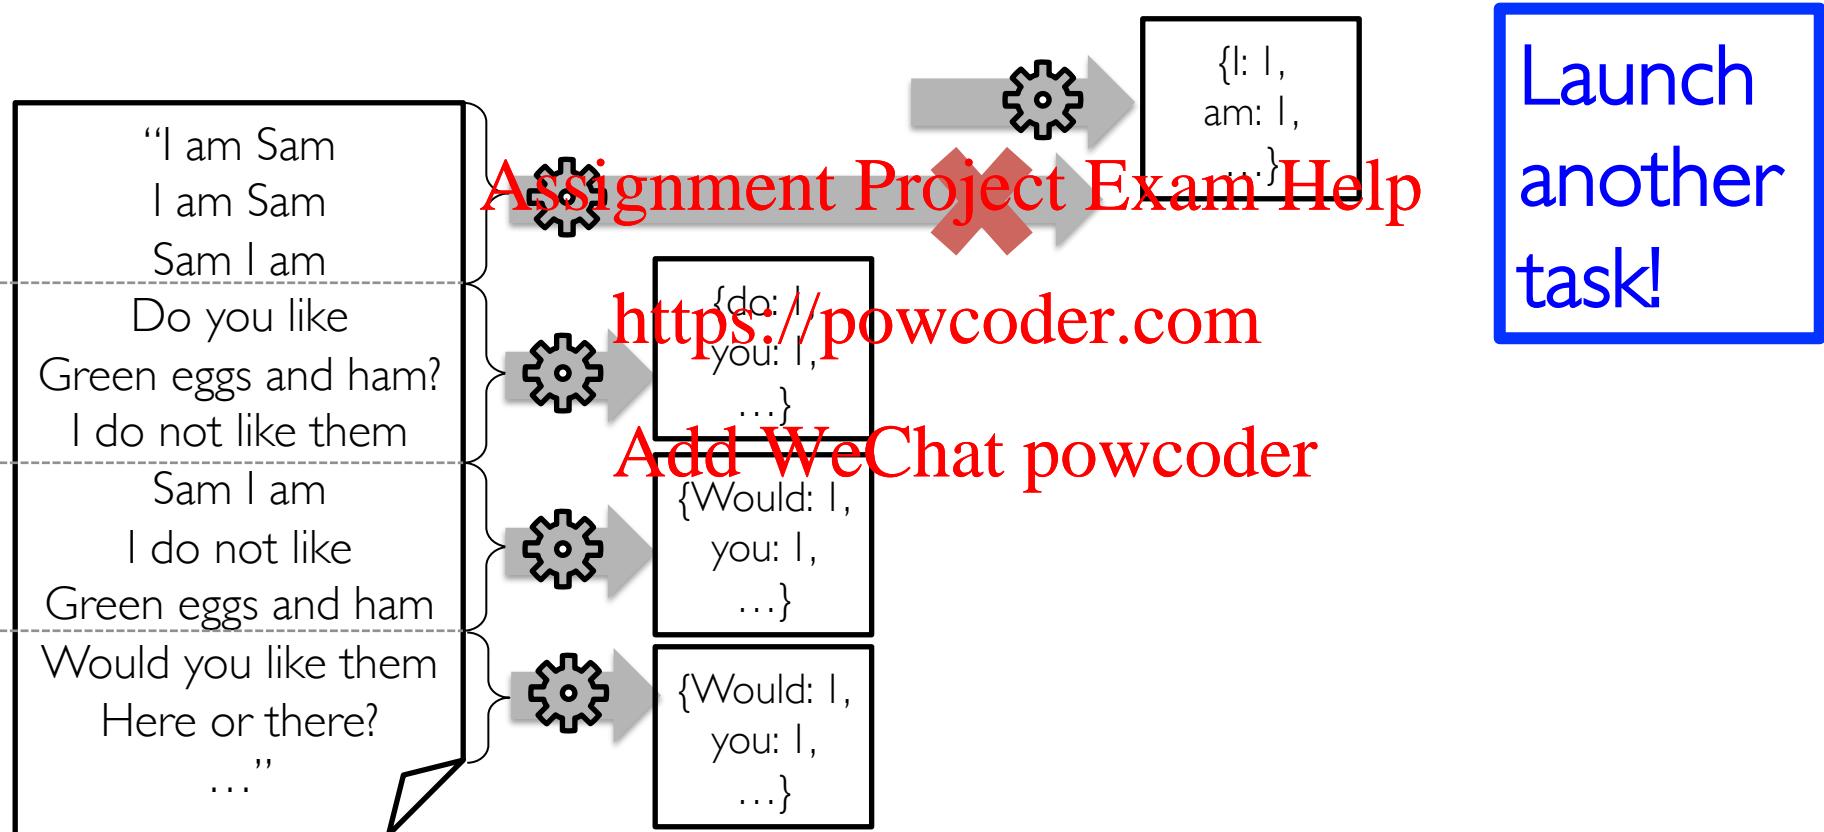


Assignment Project Exam Help

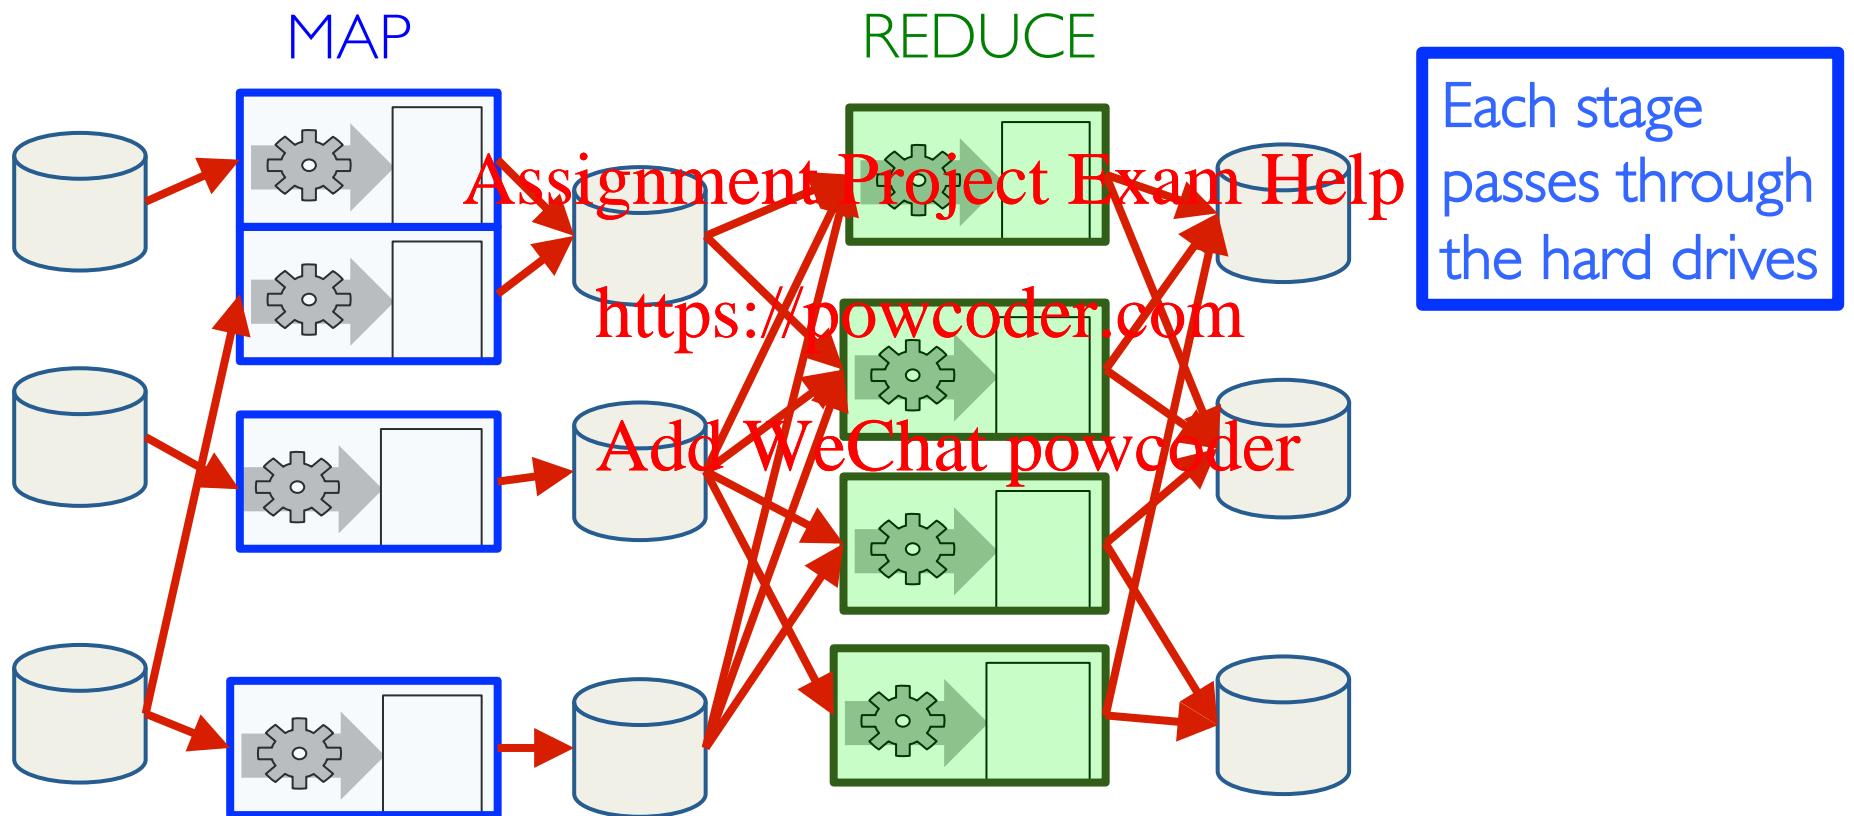
<https://powcoder.com>

Add WeChat powcoder

How Do We Deal with Slow Tasks?

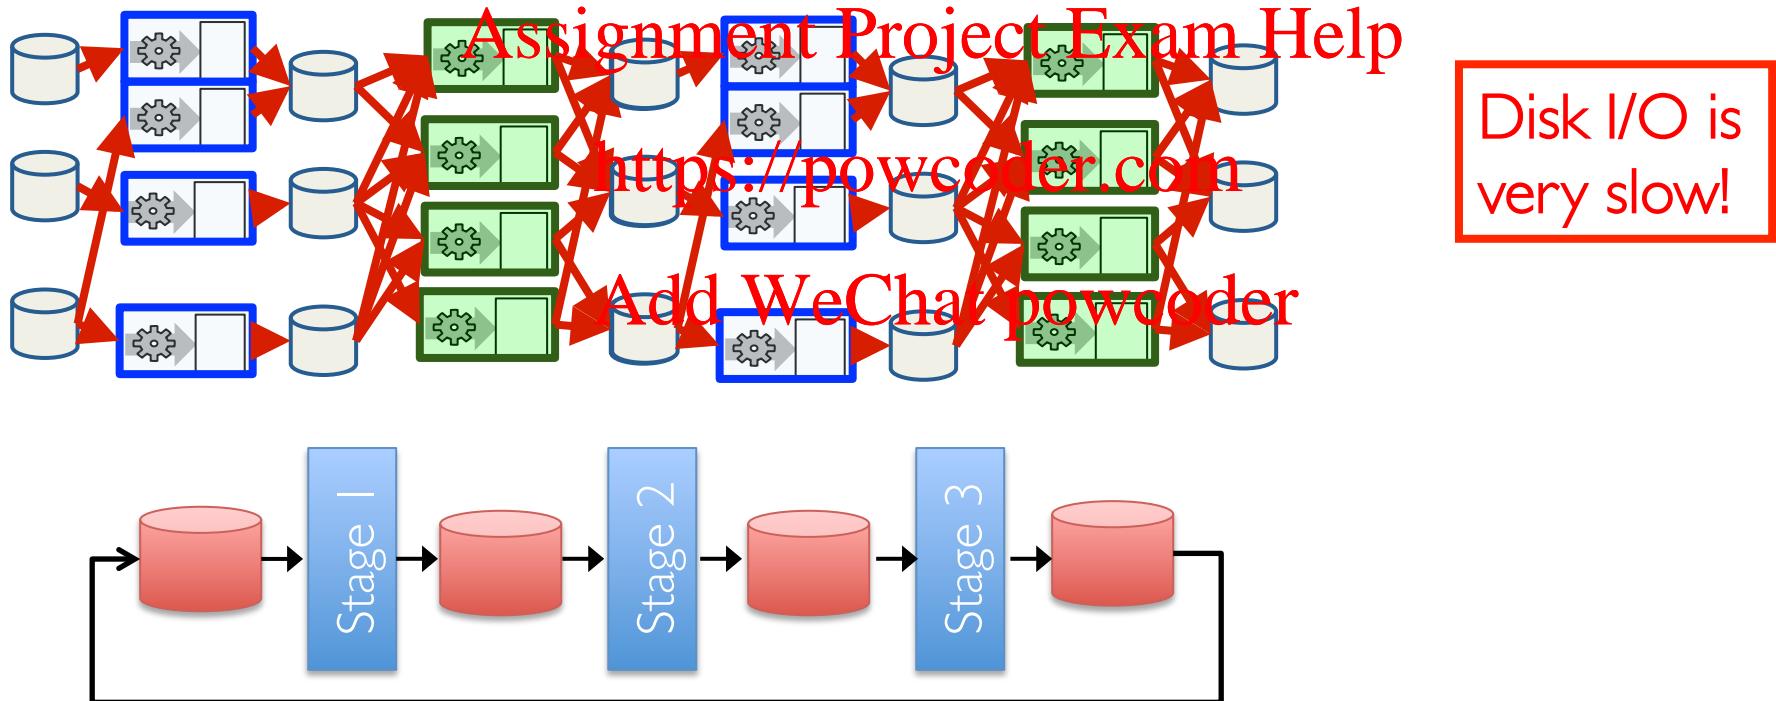


Map Reduce: Distributed Execution



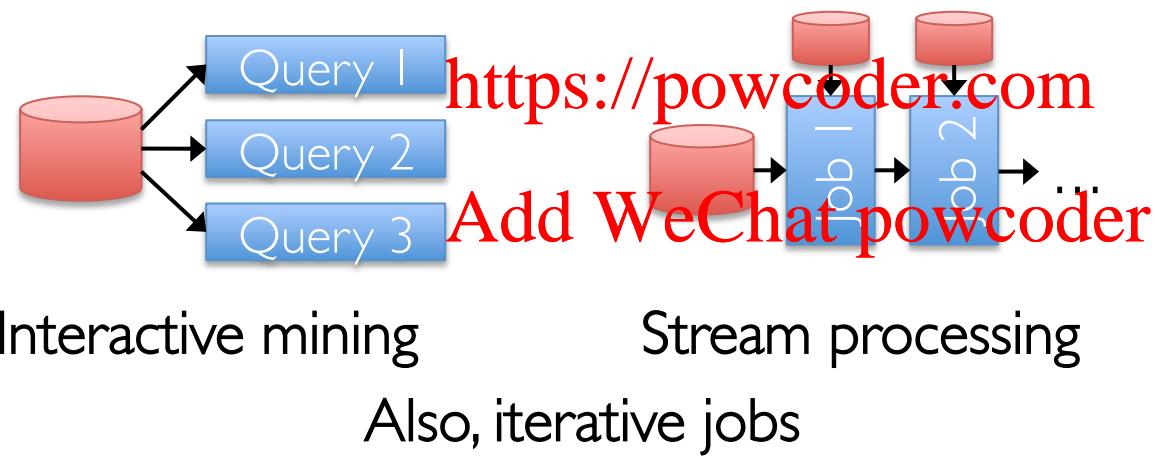
Map Reduce: Iterative Jobs

- Iterative jobs involve a lot of disk I/O for each repetition



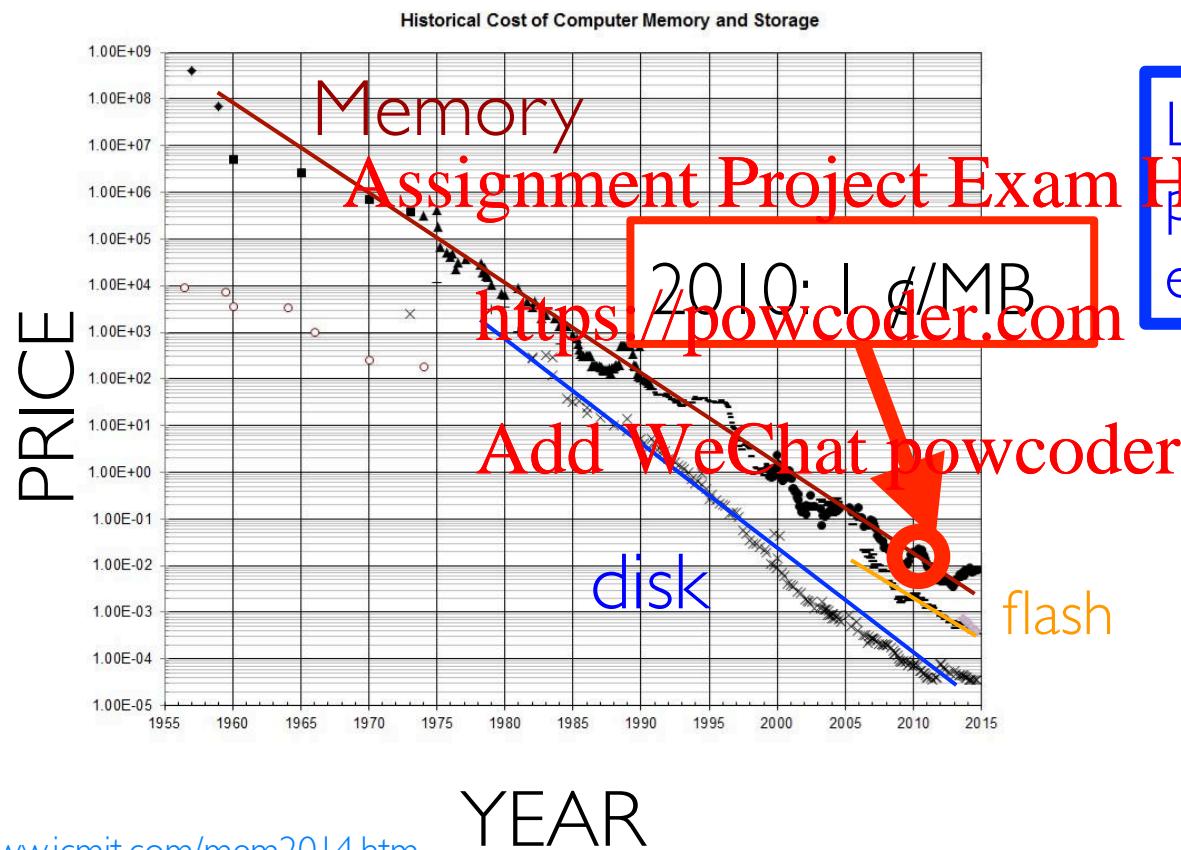
Apache Spark Motivation

- Using Map Reduce for complex jobs, interactive queries and online processing processes



Disk I/O is very slow

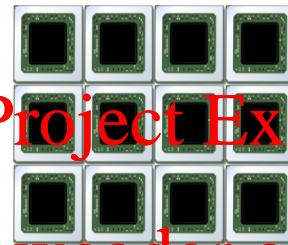
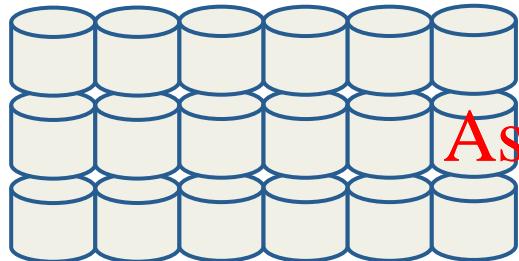
Tech Trend: Cost of Memory



Lower cost means can put more memory in each server

<http://www.jcmiit.com/mem2014.htm>

Hardware for Big Data



Assignment Project Exam Help

<https://powcoder.com>

Lots of hard drives ... and CPUs

Add WeChat powcoder



... and memory!

Opportunity

- Keep more data *in-memory*
Assignment Project Exam Help

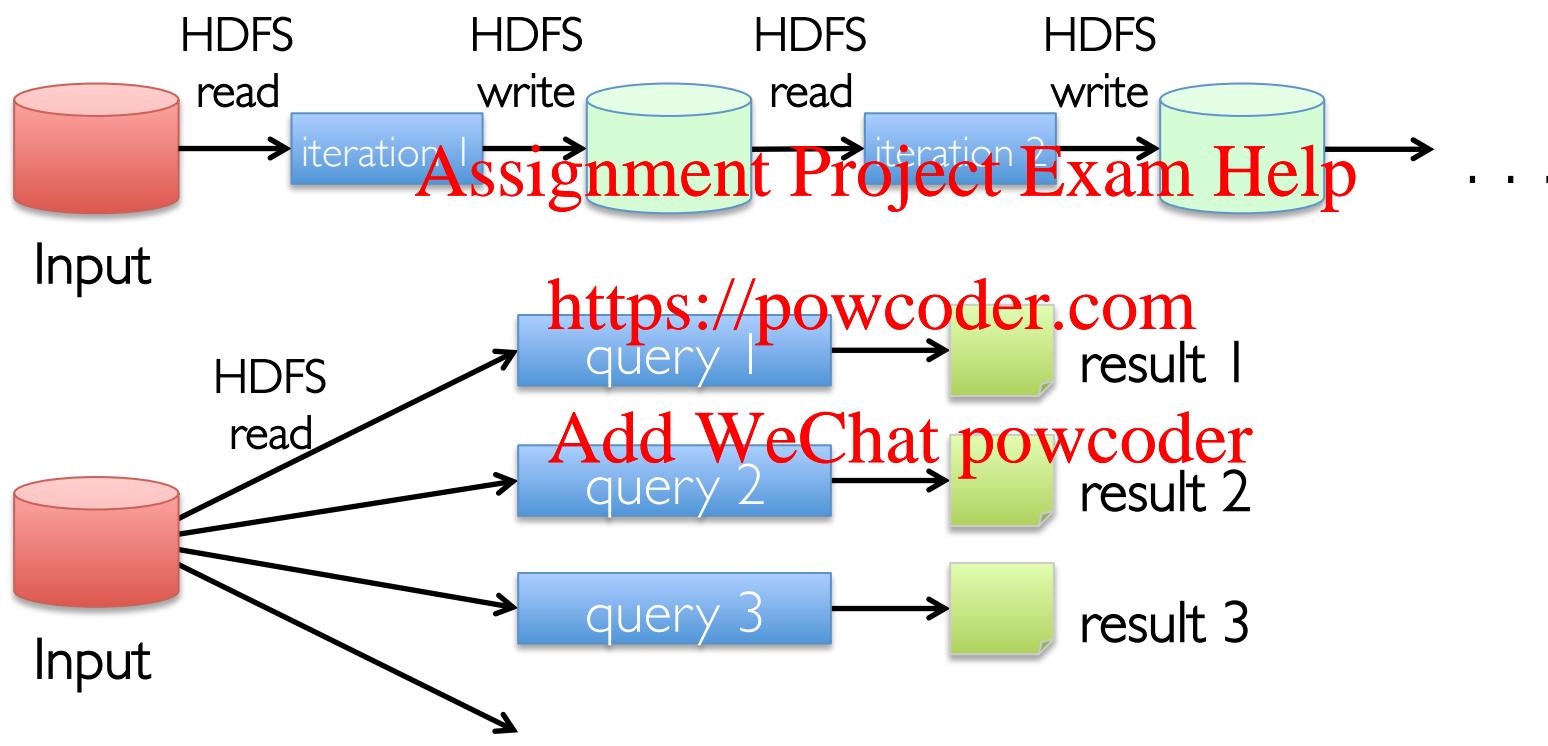
<https://powcoder.com>

- Create new distributed execution engine:
Add WeChat powcoder

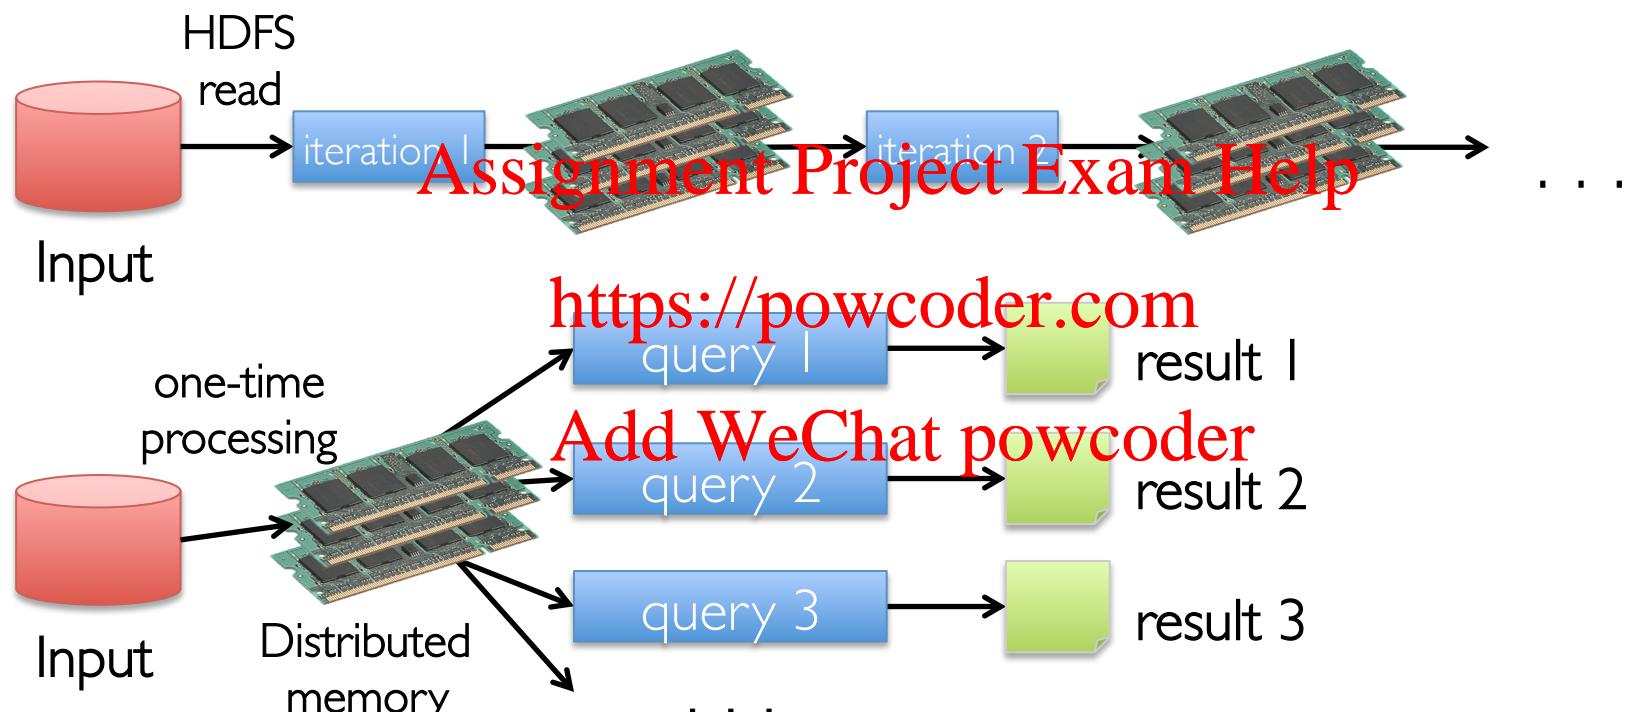


http://people.csail.mit.edu/matei/papers/2010/hotcloud_spark.pdf

Use Memory Instead of Disk



In-Memory Data Sharing



10-100x faster than network and disk

Resilient Distributed Datasets (RDDs)

- Write programs in terms of operations on distributed datasets
- Partitioned collections of objects spread across a cluster, stored in memory or on disk
 - <https://powcoder.com>
- RDDs built and manipulated through a diverse set of parallel transformations (map, filter, join) and actions (count, collect, save)
- RDDs automatically rebuilt on machine failure

The Spark Computing Framework

- Provides programming abstraction and parallel runtime to hide complexities of fault-tolerance and slow machines
<https://powcoder.com>
- “Here’s an operation, run it on all of the data”
 - » I don’t care where it runs (you schedule that)
 - » In fact, feel free to run it twice on different nodes

Spark Tools

Spark
SQL

Assignment Project Exam Help
Spark
Streaming MLlib
(machine
learning)
<https://powcoder.com>
Add WeChat powcoder

Apache Spark

Spark and Map Reduce Differences

	Hadoop Map Reduce	Spark
	Assignment Project Exam Help https://powcoder.com	
Storage	Disk only	In-memory or on disk
Operations	Map and Reduce	Map, Reduce, Join, Sampling, etc..
Execution model	Batch	Batch, interactive, streaming
Programming environments	Java	Scala, Java, R, and Python

Other Spark and Map Reduce Differences

- Generalized patterns
 ⇒ unified design for many use cases
- Lazy evaluation of the lineage graph
 ⇒ reduces wait states, better pipelining
 Add WeChat powcoder
- Lower overhead for starting jobs
- Less expensive shuffles

In-Memory Can Make a Big Difference

- Two iterative Machine Learning algorithms:



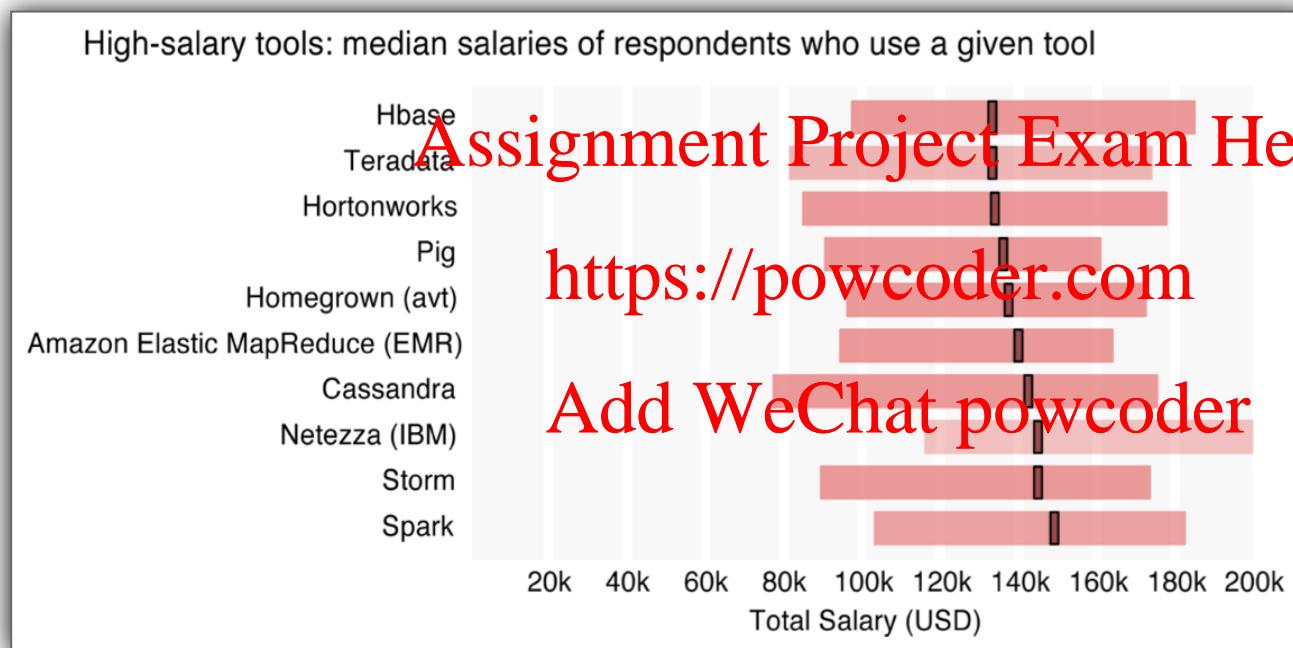
First Public Cloud Petabyte Sort

	Hadoop MR Record	Spark Record	Spark 1 PB
Data Size	102.5 TB	100 TB	1000 TB
Elapsed Time	72 mins	23 mins	234 mins
# Nodes	2100	206	190
# Cores	50400 physical	4592 virtualized	5080 virtualized
Cluster disk throughput	3150 GB/s (est.)	618 GB/s	570 GB/s
Sort Benchmark Daytona Rules	Yes	Yes	No
Network	dedicated data center, 10Gbps	virtualized (EC2) 10Gbps network	virtualized (EC2) 10Gbps network
Sort rate	1.42 TB/min	4.27 TB/min	4.27 TB/min
Sort rate/node	0.67 GB/min	20.7 GB/min	22.5 GB/min

Daytona Gray 100 TB
sort benchmark record
(tied for 1st place)

<http://databricks.com/blog/2014/11/05/spark-officially-sets-a-new-record-in-large-scale-sorting.html>

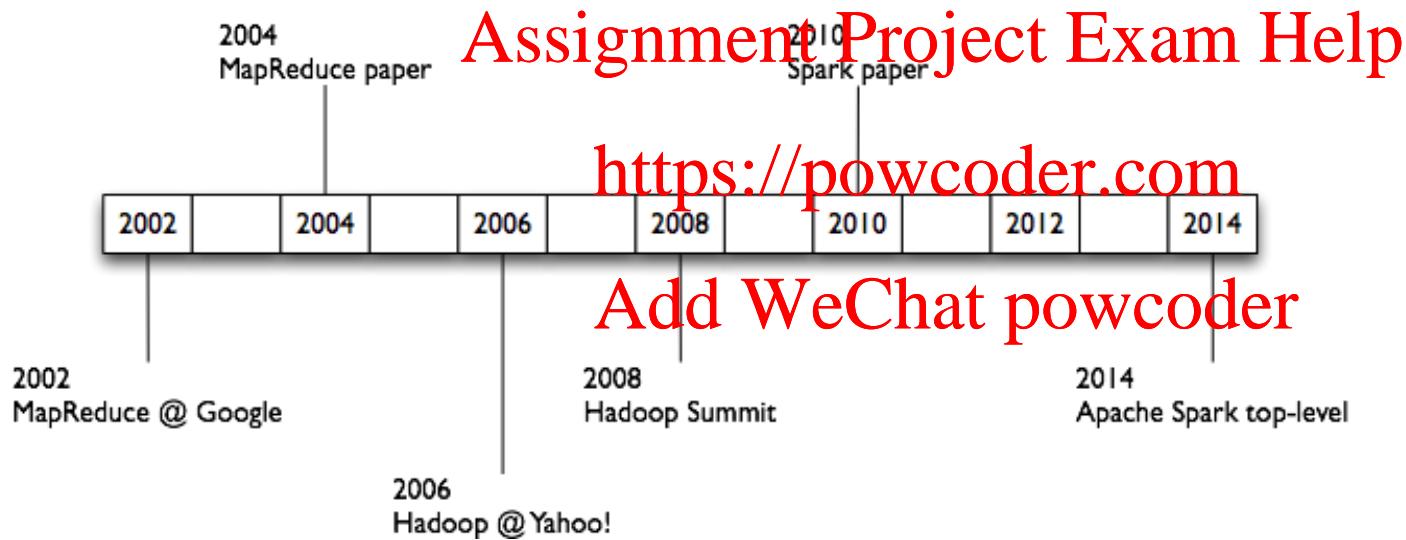
Spark Expertise Tops Big Data Median Salaries



Over 800 respondents across 53 countries and 41 U.S. states

<http://www.oreilly.com/data/free/2014-data-science-salary-survey.csp>

History Review



Historical References

- circa 1979 – Stanford, MIT, CMU, etc.: set/list operations in LISP, Prolog, etc., for parallel processing
<http://www-formal.stanford.edu/jmc/history/lisp/lisp.htm>
- circa 2004 – Google. MapReduce: Simplified Data Processing on Large Clusters
Jeffrey Dean and Sanjay Ghemawat
<http://research.google.com/archive/mapreduce.html>
- circa 2006 – Apache Hadoop, originating from the Yahoo!'s Nutch Project
Doug Cutting
<http://research.yahoo.com/files/cutting.pdf>
- circa 2008 – Yahoo!: web scale search indexing
Hadoop Summit, HUG, etc.
<http://developer.yahoo.com/hadoop/>
- circa 2009 – Amazon AWS: Elastic MapReduce
Hadoop modified for EC2/S3, plus support for Hive, Pig, Cascading, etc.
<http://aws.amazon.com/elasticmapreduce/>

Spark Research Papers

- *Spark: Cluster Computing with Working Sets*

Matei Zaharia, Mosharaf Chowdhury, Michael J. Franklin, Scott Shenker, Ion Stoica
USENIX HotCloud (2010)

Assignment Project Exam Help

people.csail.mit.edu/matei/papers/2010/hotcloud_spark.pdf

<https://powcoder.com>

- *Resilient Distributed Datasets: A Fault-Tolerant Abstraction for In-Memory Cluster Computing*

Matei Zaharia, Mosharaf Chowdhury, Tathagata Das,
Ankur Dave, Justin Ma, Murphy McCauley, Michael J. Franklin,
Scott Shenker, Ion Stoica

NSDI (2012)

usenix.org/system/files/conference/nsdi12/nsdi12-final138.pdf

Add WeChat powcoder